

## Chapter 3: Processes

肖 卿 俊

办公室：计算机楼212室

电邮： [csqjxiao@seu.edu.cn](mailto:csqjxiao@seu.edu.cn)

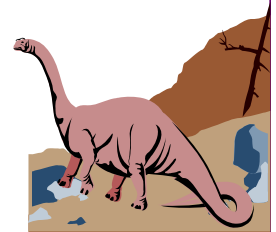
主页： <http://cse.seu.edu.cn/PersonalPage/csqjxiao>


电话： 025-52091022



# Chapter 3: Processes

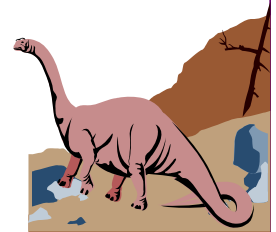
- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems






# Process Concept

- An operating system executes a variety of programs:
  - ◆ Batch system – jobs
  - ◆ Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably.






# Processes


## The Process Model

One program counter




(a)

Four program counters





(b)



(c)

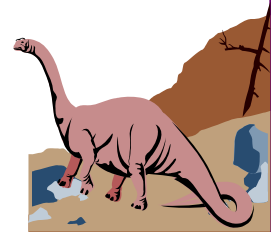
- Multiprogramming of four programs
- Conceptual model of 4 independent, sequential processes
- Only one program active at any instant






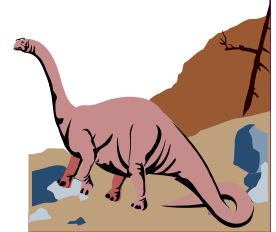

# Process Concept (Cont.)


- Process – a program in execution; process execution must progress in sequential fashion.
- A process includes:
  - ◆ program counter
  - ◆ contents of the processor's registers
  - ◆ stack
  - ◆ data section





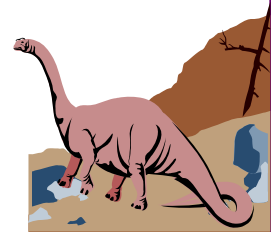
# Process in Memory







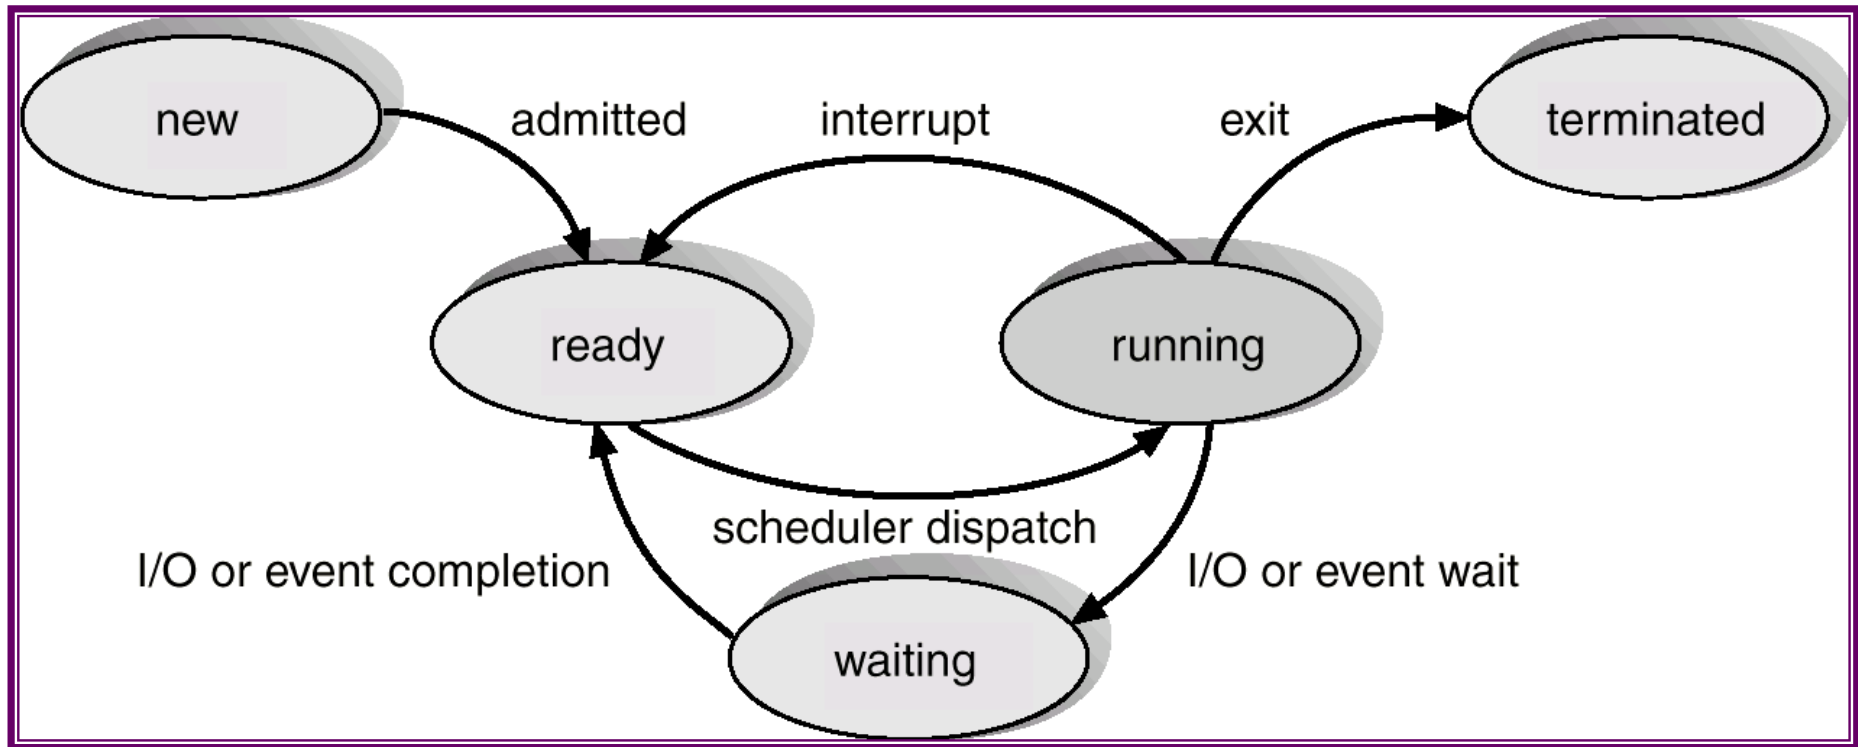
# Process State

- As a process executes, it changes *state*
  - ◆ **new**: The process is being created.
  - ◆ **running**: Instructions are being executed.
  - ◆ **waiting**: The process is waiting for some event to occur.
  - ◆ **ready**: The process is waiting to be assigned to a processor.
  - ◆ **terminated**: The process has finished execution.






# Diagram of Process State









# Process Control Block (PCB)


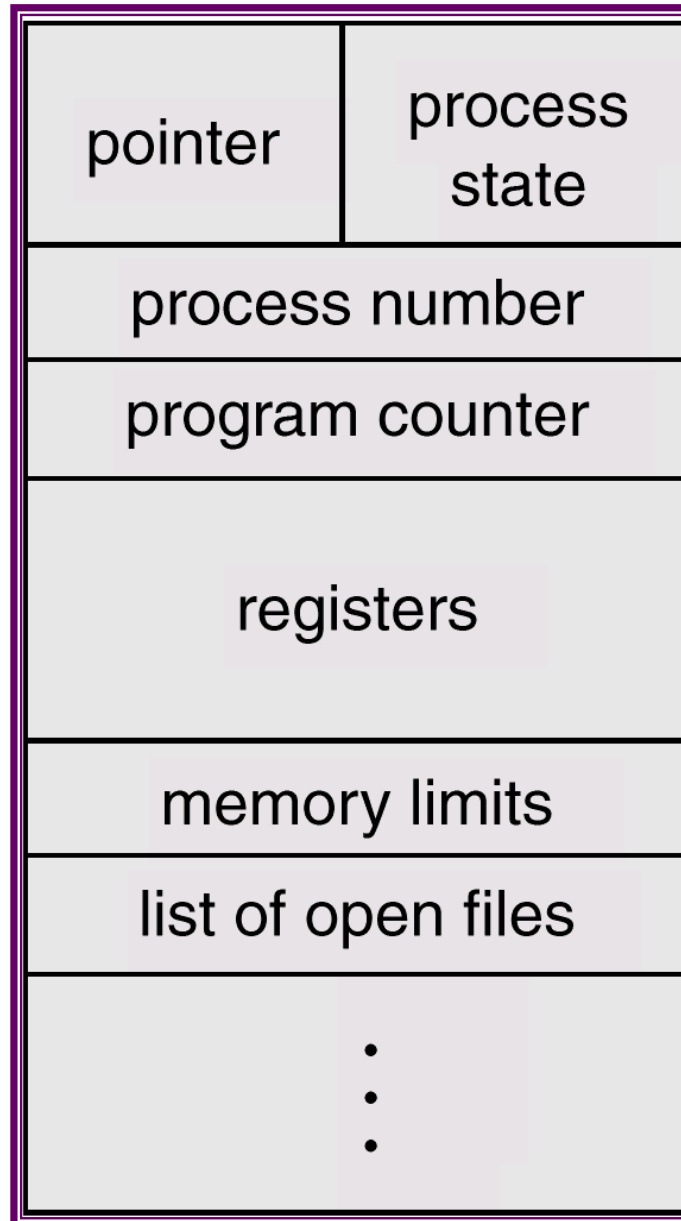
Information associated with each process.

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- File usage and I/O status information




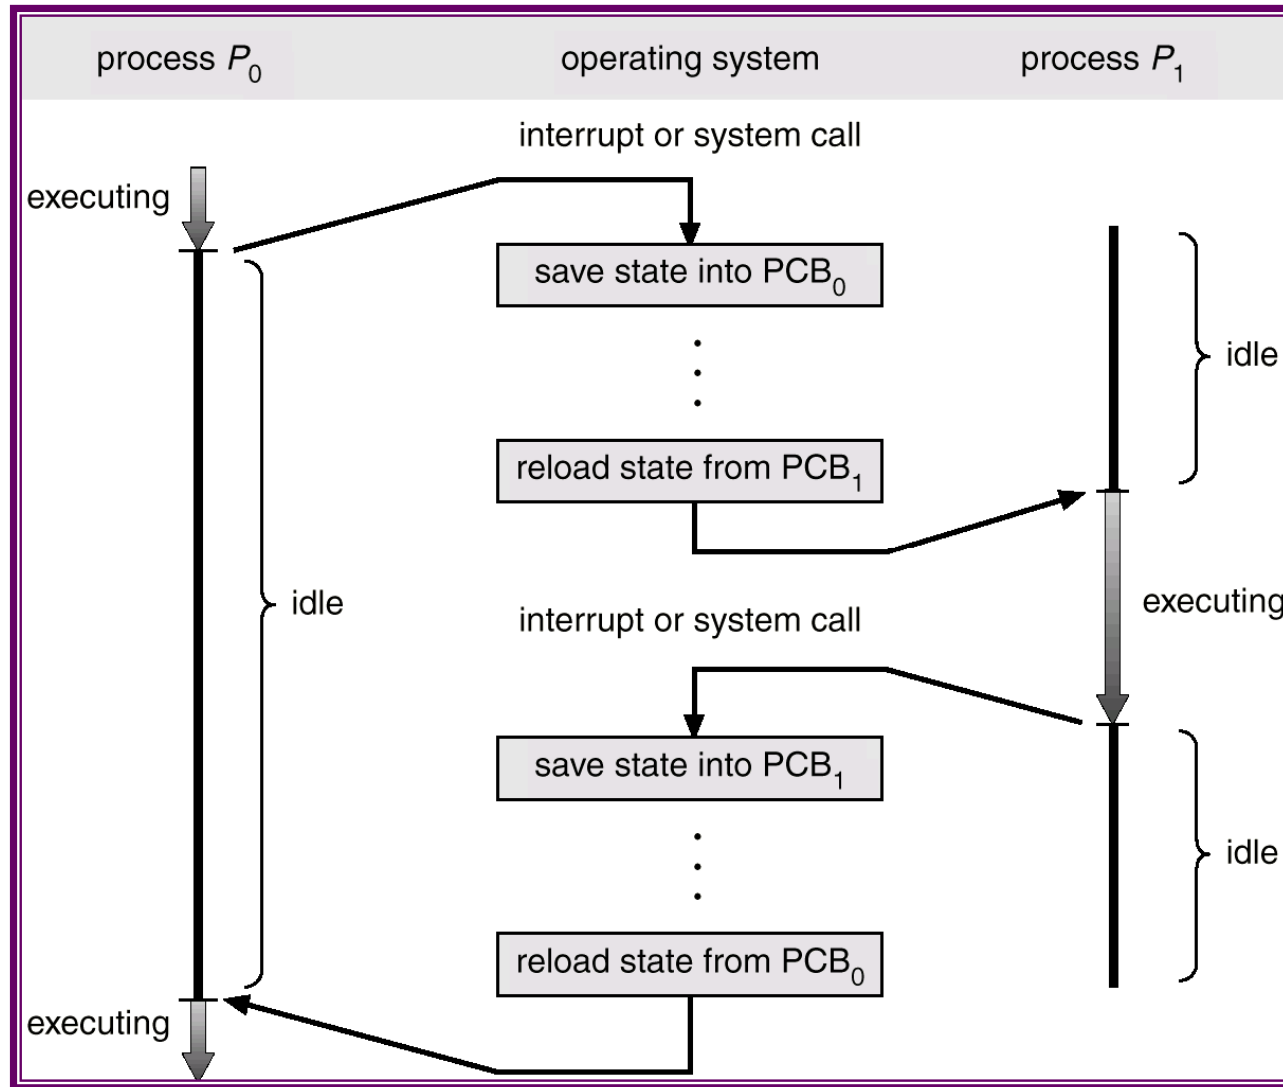


# Process Control Block (PCB)





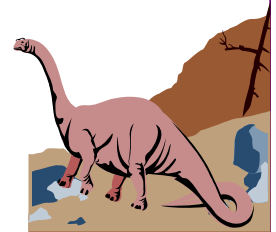
# CPU Switch From Process to Process






# Chapter 3: Processes

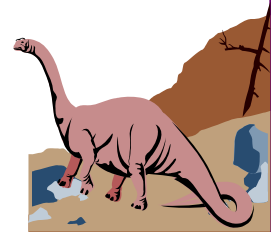
- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems







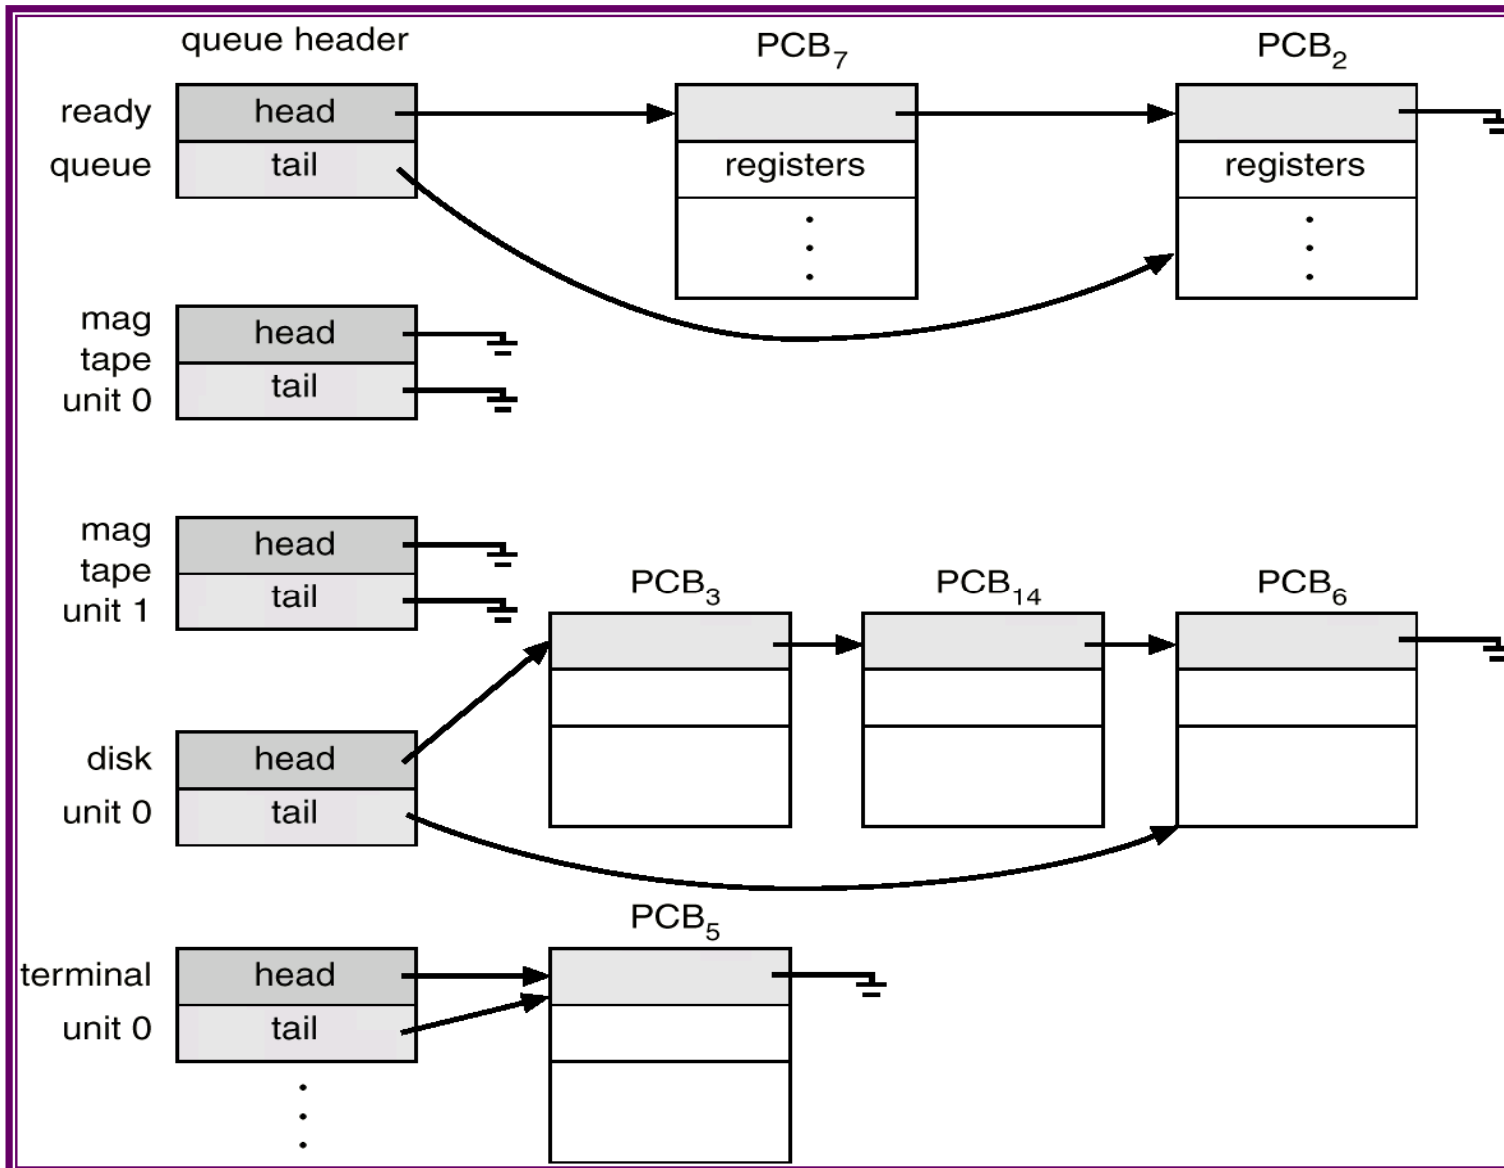
# Process Scheduling Queues


- **Job queue** – set of all processes in the system.
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute.
- **Device queues** – set of processes waiting for an I/O device.
- Process migration between the various queues.




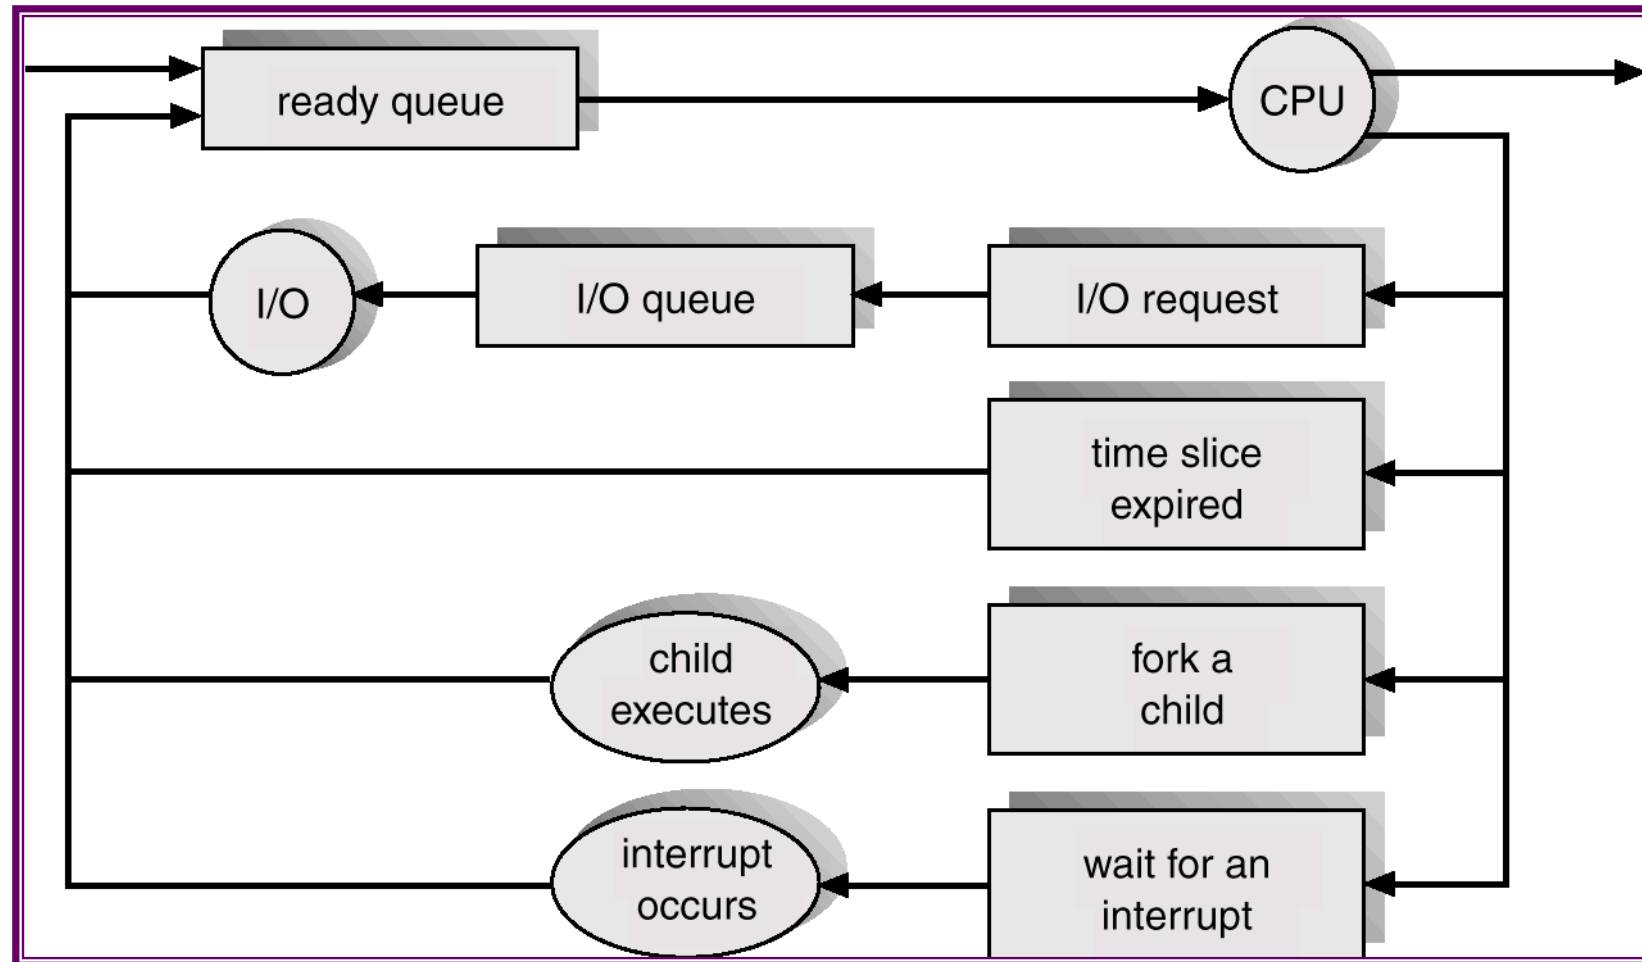



# Ready Queue And Various I/O Device Queues





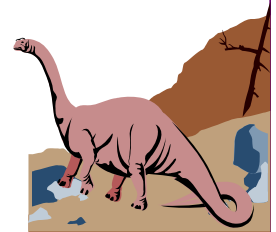
# Representation of Process Scheduling





# Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be loaded into memory for execution.
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU.








## Schedulers (Cont.)


- Short-term scheduler is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast).
- Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow).
- The \_\_\_\_\_ scheduler controls the *degree of multiprogramming*.
  - long-term
  - short-term






## Schedulers (Cont.)


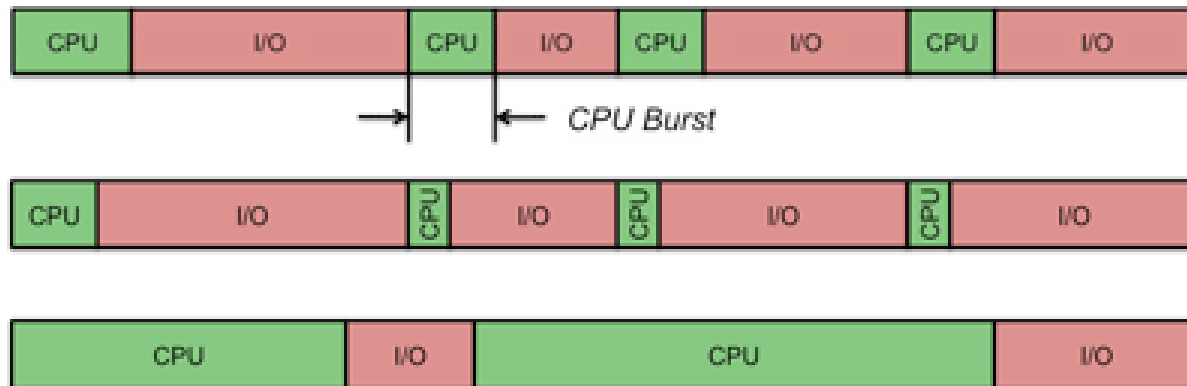
- The long-term scheduler controls the *degree of multiprogramming*.
- Long-term scheduling performs a gatekeeping function. It decides whether there's enough memory, or room, to allow new programs into the system.
- Dispatching affects processes
  - ◆ running;
  - ◆ ready;
  - ◆ blocked;
- Long-term scheduling affects processes
  - ◆ new;
  - ◆ exited






## Schedulers (Cont.)


- Processes can be described as either:
  - ◆ *I/O-bound process* – spends more time doing I/O than computations, many short CPU bursts.
  - ◆ *CPU-bound process* – spends more time doing computations; few very long CPU bursts.
- The period of computation between I/O requests is called the **CPU burst**.



# Addition of Medium-Term Scheduling



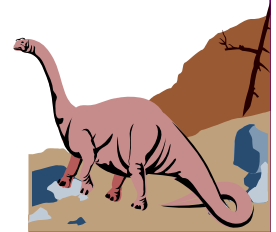
[http://en.wikipedia.org/wiki/Scheduling\\_\(computing\)#Medium-term\\_scheduling](http://en.wikipedia.org/wiki/Scheduling_(computing)#Medium-term_scheduling)




# Context Switch

## ■ What is a process context?

- ◆ The *context* of a process includes the values of CPU registers, the process state, the program counter, and other memory/file management information.

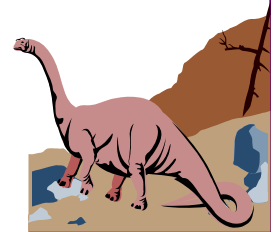





# Context Switch (Cont.)


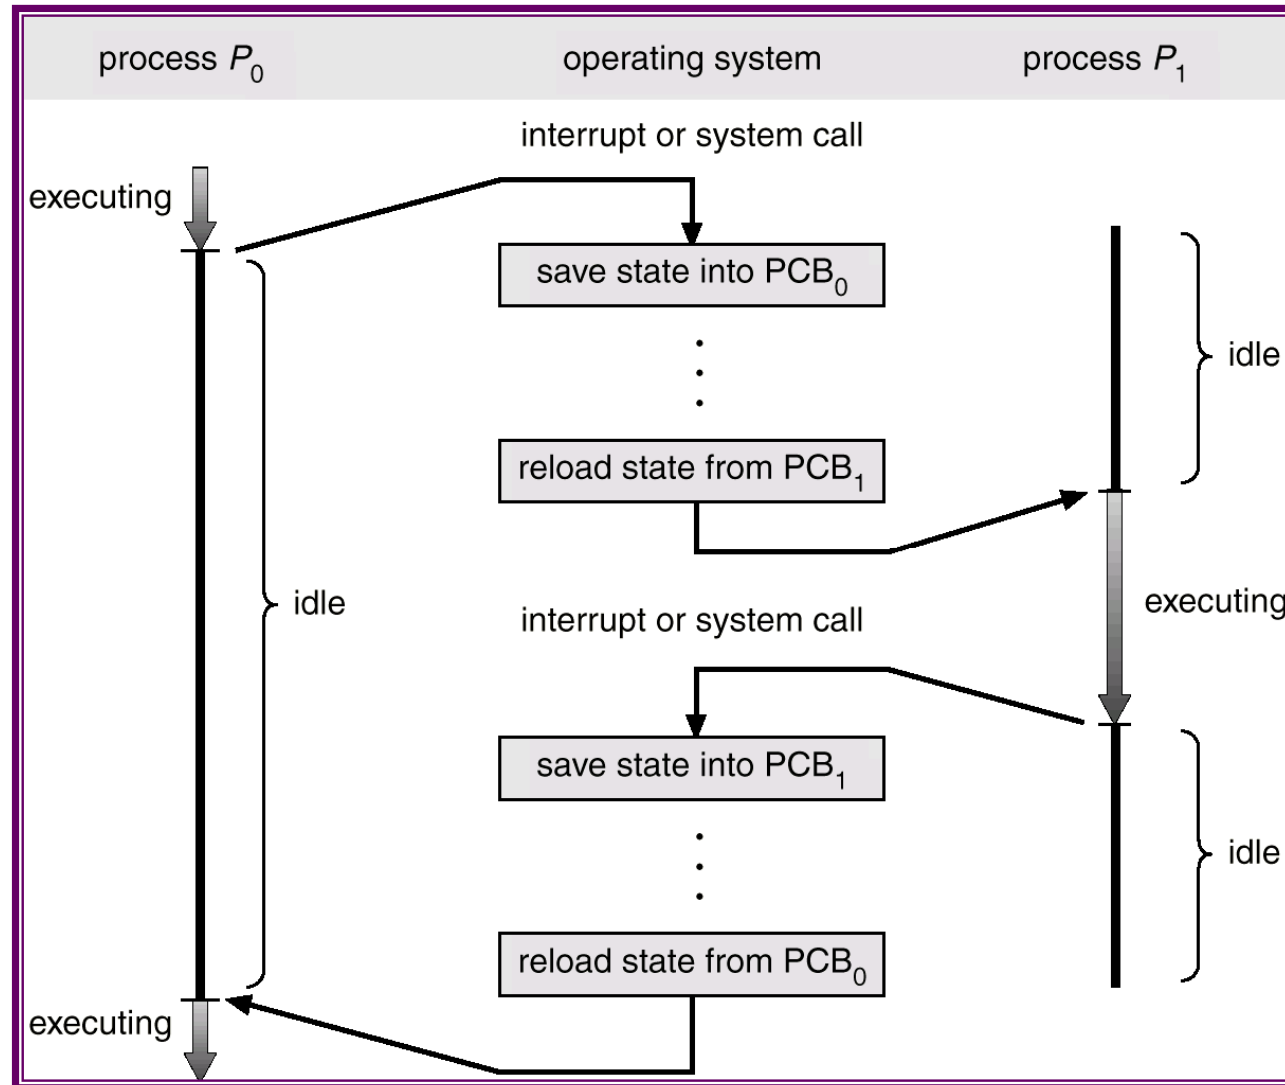
## ■ What is a context switch?


- ◆ After the CPU scheduler selects a process (from the *ready queue*) and before allocates CPU to it, the CPU scheduler must
  - ✓ save the *context* of the currently running process,
  - ✓ put it into a queue,
  - ✓ load the *context* of the selected process, and
  - ✓ let it run.





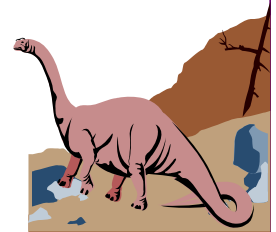
# CPU Switch From Process to Process






## Context Switch (Cont.)

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support.

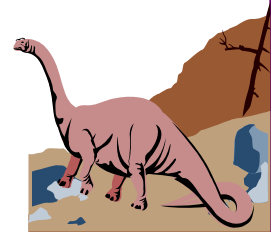







# Chapter 3: Processes

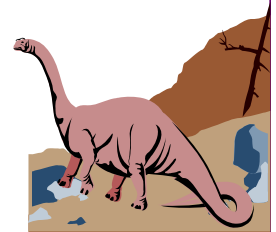
- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems






# Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes.
- Resource sharing
  - ◆ Parent and children share all resources.
  - ◆ Children share subset of parent's resources.
  - ◆ Parent and child share no resources.
- Execution
  - ◆ Parent and children execute concurrently.
  - ◆ Parent waits until children terminate.





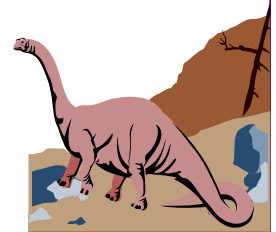
# Process Creation (Cont.)

## ■ Address space

- ◆ Child duplicate of parent.
- ◆ Child has a program loaded into it.



## ■ UNIX examples

- ◆ **fork** system call creates new process
- ◆ **exec** system call used after a **fork** to replace the process' memory space with a new program.



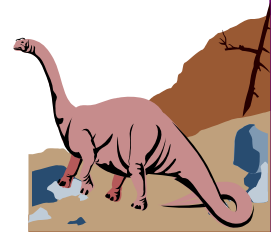


# Process Creation (UNIX)


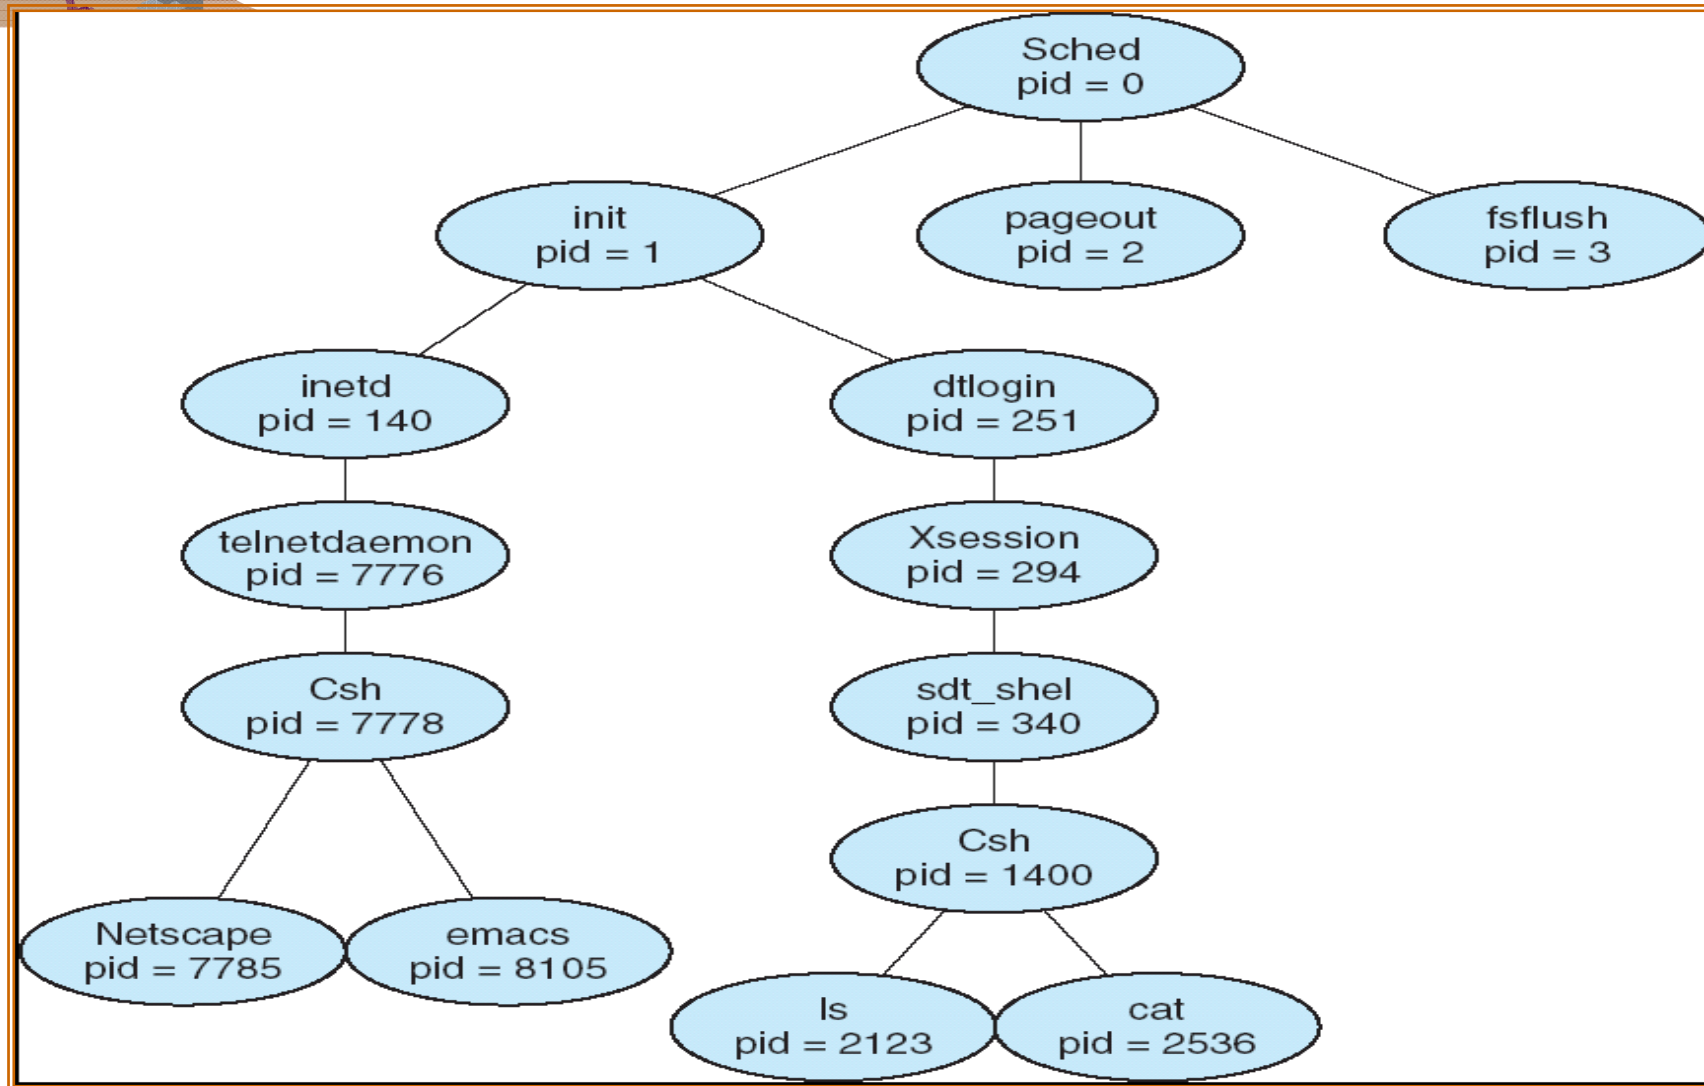




```
pid = fork();  
if (pid < 0){/*error occured*/  
    fprintf(stderr, "Fork failed");  
    exit(-1);}  
else if(pid == 0){/*child process*/  
    execlp("/bin/ls","ls",NULL);  
}  
else {/*parent process*/  
    wait(NULL);  
    printf("Child Complete");  
    exit(0);  
}
```




# Processes Tree on Solaris






# Process Termination

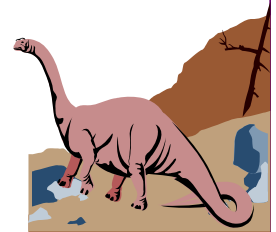
- Process executes last statement and asks the operating system to delete it (**exit**).
  - ◆ Output data from child to parent (via **wait**).
  - ◆ Process' resources are deallocated by OS.
- Parent may terminate execution of children processes (**abort**).
  - ◆ Child has exceeded allocated resources.
  - ◆ Task assigned to child is no longer required.
  - ◆ Parent is exiting.
    - ✓ Operating system does not allow child to continue if its parent terminates.
    - ✓ Cascading termination.Z






# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Communication in Client-Server Systems

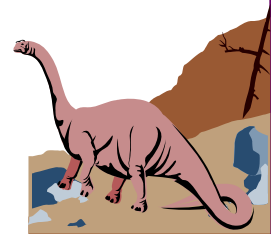







# Cooperating Processes

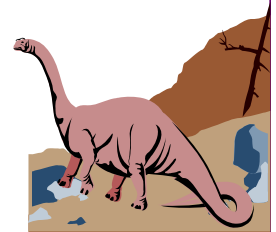
- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - ◆ Information sharing
  - ◆ Computation speed-up
  - ◆ Modularity
  - ◆ Convenience






# Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process.
  - ◆ *unbounded-buffer* places no practical limit on the size of the buffer.
  - ◆ *bounded-buffer* assumes that there is a fixed buffer size.





# Bounded-Buffer – Shared-Memory Solution

## ■ Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

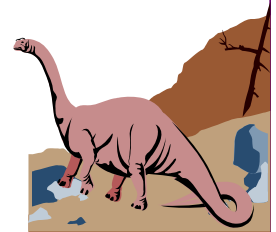
```
    ...
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

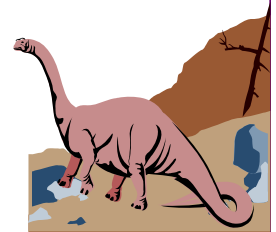




# Bounded-Buffer – Producer Process

```
item nextProduced;
```

```
while (1) {  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

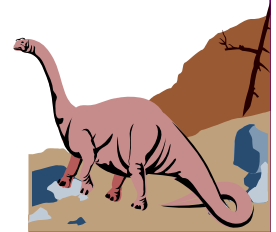





# Bounded-Buffer – Consumer Process

```
item nextConsumed;
```

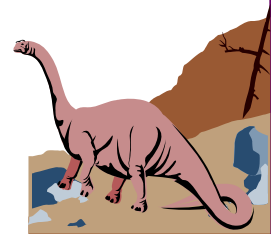
```
while (1) {  
    while (in == out)  
        ; /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
}
```





# Bounded-Buffer


- Solution is correct, but can only use `BUFFER_SIZE-1` elements






# Interprocess Communication (IPC)

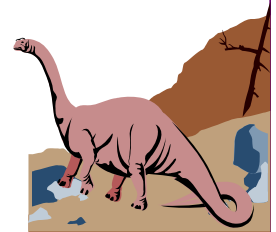
- Mechanism for processes to communicate and to synchronize their actions.
- Message system – processes communicate with each other without resorting to shared variables.






# Interprocess Communication (Cont.)

- IPC facility provides two operations:
  - ◆ **send**(*message*) – message size fixed or variable
  - ◆ **receive**(*message*)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - ◆ establish a *communication link* between them
  - ◆ exchange messages via send/receive
- Implementation of communication link
  - ◆ physical (e.g., shared memory, hardware bus)
  - ◆ logical (e.g., logical properties)

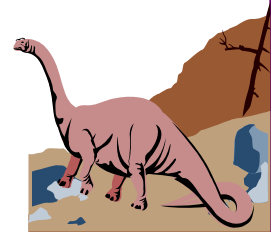







# Implementation Questions

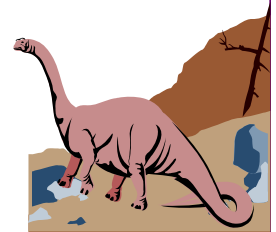
- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?






# Direct Communication

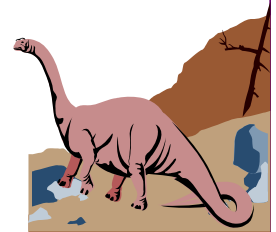
- Processes must name each other explicitly:
  - ◆ **send** ( $P$ , *message*) – send a message to process  $P$
  - ◆ **receive**( $Q$ , *message*) – receive a message from process  $Q$
- Properties of communication link
  - ◆ Links are established automatically.
  - ◆ A link is associated with exactly one pair of communicating processes.
  - ◆ Between each pair there exists exactly one link.
  - ◆ The link may be unidirectional, but is usually bi-directional.






# Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports).
  - ◆ Each mailbox has a unique id.
  - ◆ can communicate only if they share a mailbox.
- Properties of communication link
  - ◆ Link established only if processes share a common mailbox
  - ◆ A link may be associated with many processes.
  - ◆ Each pair of processes may share several communication links.
  - ◆ Link may be unidirectional or bi-directional.





# Indirect Communication

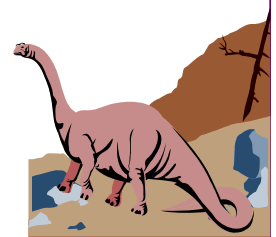
## ■ Operations


- ◆ create a new mailbox
- ◆ send and receive messages through mailbox
- ◆ destroy a mailbox

## ■ Primitives are defined as:

**send**(*A, message*) – send a message to mailbox *A*

**receive**(*A, message*) – receive a message from mailbox *A*






# Indirect Communication


## ■ Mailbox sharing

- ◆  $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A.
- ◆  $P_1$  sends;  $P_2$  and  $P_3$  receive.
- ◆ Who gets the message?

## ■ Solutions

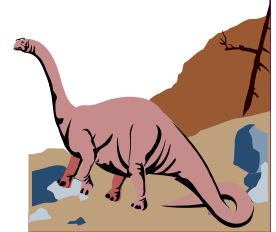
- ◆ Allow a link to be associated with at most two processes.
- ◆ Allow only one process at a time to execute a receive operation.
- ◆ Allow the system to select arbitrarily the receiver.  
Sender is notified who the receiver was.






# Synchronization

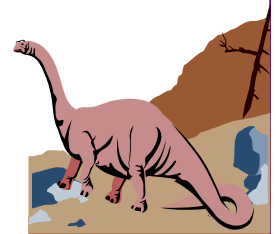
- Message passing may be either blocking or non-blocking.
- **Blocking** is considered **synchronous**
- **Non-blocking** is considered **asynchronous**
- **send** and **receive** primitives may be either blocking or non-blocking.






# Buffering

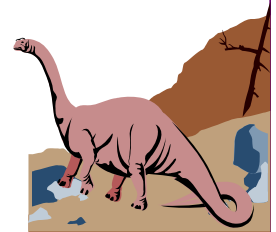
- Queue of messages attached to the link; implemented in one of three ways.
  1. Zero capacity – 0 messages  
Sender must wait for receiver (rendezvous).
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full.
  3. Unbounded capacity – infinite length  
Sender never blocks.






# Chapter 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication
- Communication in Client-Server Systems

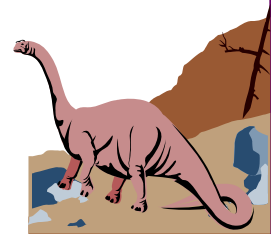







# Client-Server Communication

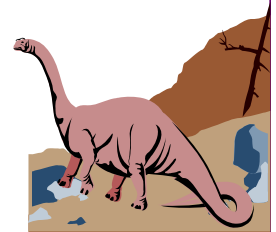
- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)







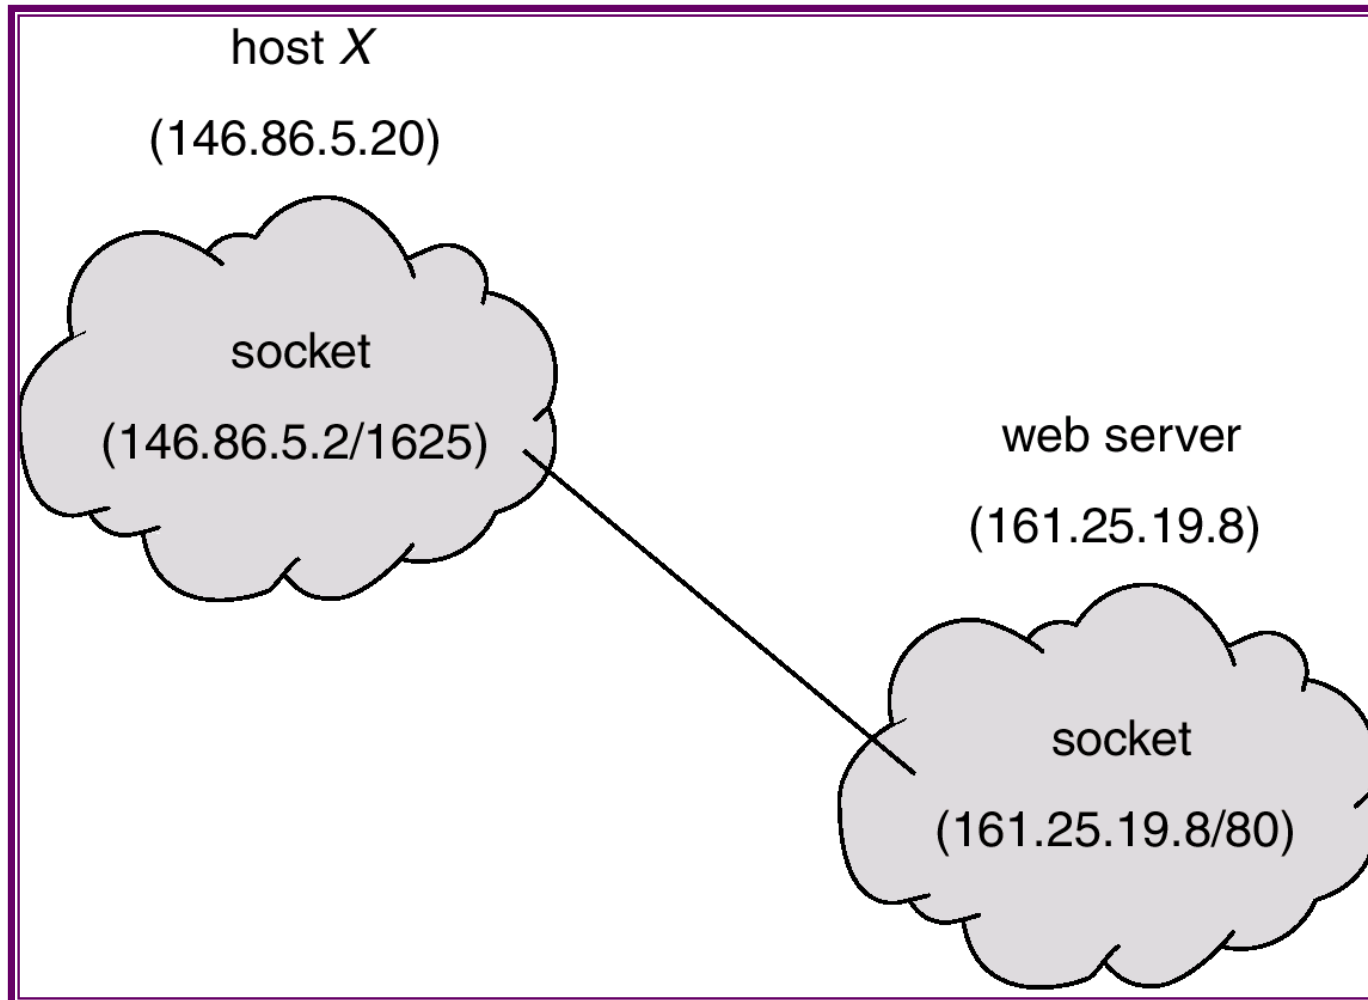
# Sockets


- A socket is defined as an *endpoint for communication*.
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets.





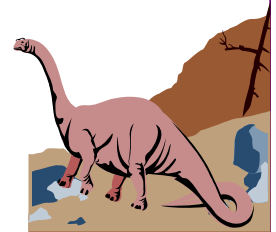
# Socket Communication







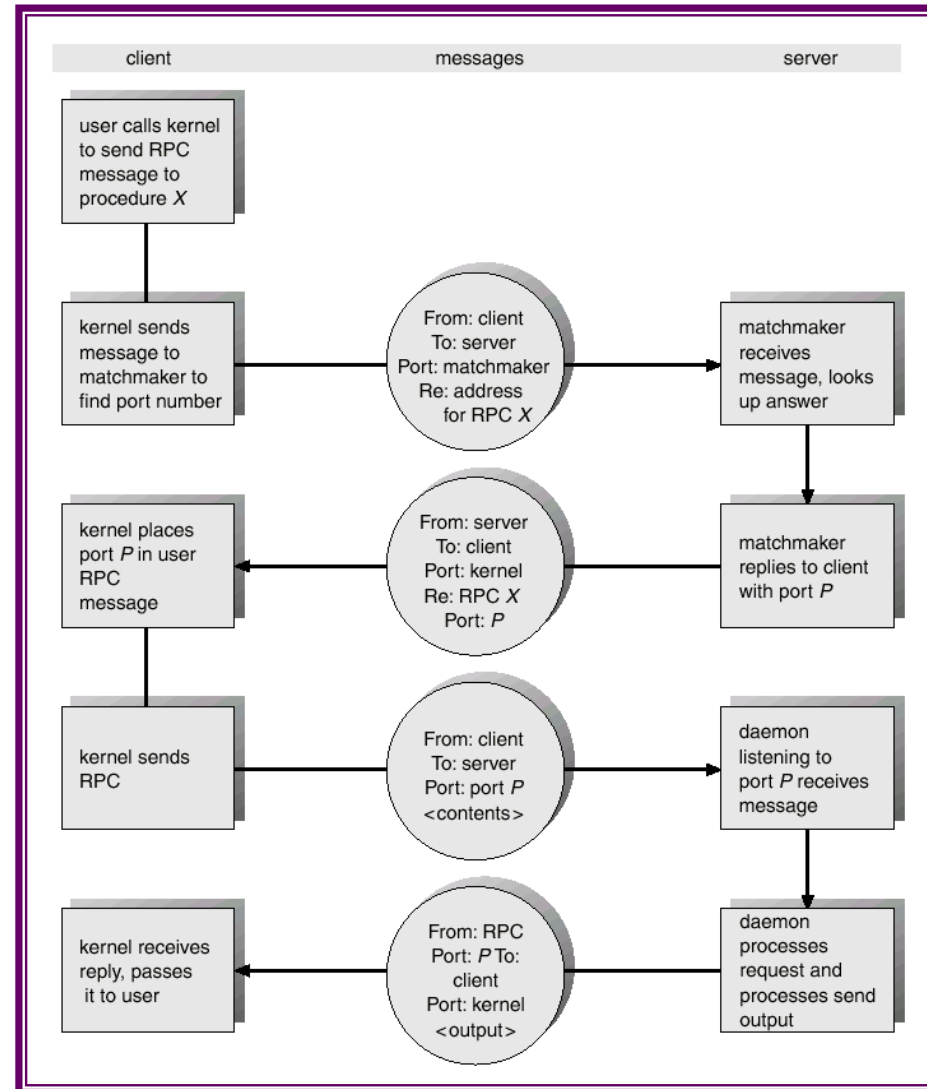
# Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshalls* the parameters.
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.







# Execution of RPC






# Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- RMI allows a Java program on one machine to invoke a method on a remote object.





# Marshalling Parameters

